

# *MCPProf*: Memory and Communication Profiler

Imran Ashraf  
Computer Engineering Lab, TU Delft, The Netherlands  
I.Ashraf@tudelft.nl

April 27, 2017

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>Licensing</b>	<b>2</b>
<b>3</b>	<b>Reference <i>MCPProf</i></b>	<b>2</b>
3.1	To cite MCPProf . . . . .	2
<b>4</b>	<b>Supported Platforms</b>	<b>2</b>
<b>5</b>	<b>Availability</b>	<b>3</b>
<b>6</b>	<b>Required Packages</b>	<b>3</b>
<b>7</b>	<b>Installation</b>	<b>3</b>
<b>8</b>	<b>Usage</b>	<b>4</b>
8.1	Profiling Given Tests . . . . .	4
8.2	Profiling Your Own Example . . . . .	6
<b>9</b>	<b><i>MCPProf</i> Input Options</b>	<b>6</b>
<b>10</b>	<b><i>MCPProf</i> Generated Output</b>	<b>7</b>
10.1	Engine 1 Output . . . . .	8
10.2	Engine 2 Output . . . . .	11
10.3	Engine 3 Output . . . . .	11
<b>11</b>	<b>Frequently Encountered Problems</b>	<b>13</b>
11.1	Pin: No such file or directory . . . . .	13
11.2	Pin Injection Mode Error . . . . .	13
<b>12</b>	<b>Contact</b>	<b>13</b>

# 1 Introduction

*MCPProf* is a memory and communication profiler. It traces memory read/writes and reports memory accesses by various functions in the application as well as the data-communication between functions. The information is obtained by performing dynamic binary instrumentation by utilizing Intel Pin [1, 2] framework. This manual explains the process of setting up *MCPProf* and using it.

## 2 Licensing

```
Copyright (c) 2014-2015 TU Delft, The Netherlands.
All rights reserved.

MCPProf is free software: you can redistribute it and/or modify it under the
terms of the GNU Lesser General Public License as published by the Free
Software Foundation, either version 3 of the License, or (at your option)
any later version.

MCPProf is distributed in the hope that it will be useful, but WITHOUT ANY
WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR
A PARTICULAR PURPOSE. See the GNU Lesser General Public License for more details.

You should have received a copy of the GNU Lesser General Public License along
with MCPProf. If not, see <http://www.gnu.org/licenses/>.
```

## 3 Reference *MCPProf*

### 3.1 To cite MCPProf

```
@phdthesis{phdthesisAshraf,
author = "I. Ashraf",
title = "Communication Driven Mapping of Applications on Multicore Platforms",
school = "Delft University of Technology",
address = "Delft, Netherlands",
month = "April",
year = "2016"
}

@inproceedings{mcpProf,
author = "I. Ashraf and V.M. Sima and K.L.M. Bertels",
title = "Intra-Application Data-Communication Characterization",
booktitle = "Proc. 1st International Workshop on Communication Architectures at Extreme Scale",
address = "Frankfurt, Germany",
month = "July",
year = "2015",
}
```

## 4 Supported Platforms

*MCPProf* relies on Intel Pin, so it can be used on any 64-bit Linux platform for which Pin is available. We have used *MCPProf* on 64-bit Ubuntu 12.04, Ubuntu 14.04, Ubuntu 16.04 on real machines as well as virtual machines running in virtualbox.

## 5 Availability

*MCPProf* can be downloaded from [3].

You can also clone *MCPProf* from the git repository by:

```
git clone https://imranashraf@bitbucket.org/imranashraf/mcprof.git
```

In this way, you don't need to download *MCPProf* each time a new version is released. You can update your *MCPProf* by going inside the directory where you have cloned *MCPProf* and executing the following commands:

```
make clean
git fetch
git pull origin master
```

## 6 Required Packages

In order to setup and use *MCPProf* the following two packages are required:

- Intel Pin DBI framework [2] Revision 62732 or higher
- g++ compiler with support for C++11X
- libelf-dev
- libdwarf-dev
- graphviz Dot utility for converting the generated communication graphs from DOT to pdf formats
- gnuplot to plot communication matrix as graph

Intel Pin does not need installation. You just need to download the archive and extract to some location. For the remaining packages, you can search on internet how to install these packages in your linux distribution. In Ubuntu 14.04, these packages can be installed, for example, by the following commands:

```
sudo apt-get install g++
sudo apt-get install libelf-dev
sudo apt-get install libdwarf-dev
sudo apt-get install graphviz
sudo apt-get install gnuplot
```

## 7 Installation

*MCPProf* uses Makefile to compile the sources. In order to compile *MCPProf* from sources on Linux, the following steps can be performed.

- Download Pin and copy and extract it to the directory where you want to keep Pin.
- Define a variable `PIN_ROOT` by running the following commands:

```
export PIN_ROOT=/<absolute path to pin directory>
```

- Add Pin to your path by the following command:

```
export PATH=$PIN_ROOT:$PATH
```

- You can also add these lines, for instance, to your `.bashrc` in case you are using `bash` to export these variables automatically on opening a terminal.
- Download *MCPProf* and copy and extract it to the directory where you want to compile it.
- Go the *MCPProf* directory and run the following command to compile it:

```
make
```

If every thing goes fine, you will see a directory `obj-intel64` (or `obj-ia32` depending upon your architecture). This directory will contain the executables and object files generated as a result of the compilation. The important files are:

- **mcprof.so** which is the tool. This will be used to profile the applications as explained in Section 8.
- executable files of the test applications available in **tests** directory. These executables can be used as test inputs.

## 8 Usage

In order to explain the usage of *MCPProf* we will use the example application listed in Figure 1. The complete source-code is available in **tests** directory of source package. In this application, 4 `int` arrays are created on source lines 23, 24, 25 and 26. These arrays are initialized in `initVecs` function. The sum and difference of the elements of these arrays are computed in `sumVecs` and `diffVecs` functions, respectively. Finally, these arrays are free on lines 32-35.

### 8.1 Profiling Given Tests

This example will be compiled during the default compilation of the *MCPProf* discussed in Section 7. In order to profile this application by *MCPProf* you can give the following command:

---

```

1 void test(int count){
2     printf("Call number %d to test()\n",count);
3     if(count)    test(count-1);
4     else        return;
5 }
6 void initVecs() {
7     for(int i = 0; i < nElem; i++)
8     {
9         srcArr1[i]=i*5 + 7;
10        srcArr2[i]=2*i - 3;
11    }
12 }
13 void sumVecs() {
14     initVecs();
15     for(int i = 0; i < nElem; i++)
16     {
17         sumArr[i] = srcArr1[i] + coeff * srcArr2[i];
18     }
19     initVecs();
20 }
21 void diffVecs() {
22     test(3);
23     for(int i = 0; i < nElem; i++) {
24         diffArr[i] = coeff * (srcArr1[i] - srcArr2[i]);
25     }
26 }
27 void process() {
28     initVecs();
29     sumVecs();
30     diffVecs();
31 }
32 int main() {
33     nElem = 5000;
34     nBytes = nElem*sizeof(TYPE);
35     srcArr1 = malloc(nBytes); // similarly other allocations
36
37     int i;
38     for(i=0; i<3; i++)
39         process();
40
41     printf("output : %d\n",sumArr[1]+diffArr[2]);
42
43     free(srcArr1); // similarly other frees
44     return 0;
45 }

```

---

Figure 1: Example of an application processing some arrays.

```
make vectOps.test
```

Similarly, other tests can also be executed by replacing the `<vectOps>.test` with the `<test application name>.test` as given in `tests` directory. This will generate the output information depending upon the selected engine. The details of the generated output are provided in Section 10.

## 8.2 Profiling Your Own Example

In order to provide an example of how you can compile and profile your own application, the same `vectOps` example is provided in directory `yourApp`. You can copy this directory to any location you like your Linux machine. In order to profile this application, a `makefile` is provided in this directory containing all the rules to compile and profile this application. You need to make some changes before profiling this application.

- Modify the path to Pin directory on line 1 in the makefile.
- Modify the path to *MCPProf* directory on line 2 in the makefile.

In order to avoid these changes in each makefile, you can also add these paths in `.bashrc` file as:

```
export PIN_ROOT=/<absolute path to pin directory>
export PATH=$PIN_ROOT:$PATH
export MCPROF_ROOT=/<absolute path to mcprof directory>
```

It should be noted that the *MCPProf* options are provided in `MCPROF_OPT` variable in `makefile`. You can modify these options as required. The details of these options are available in Section 9. Once these modifications are performed, this application can be compiled by:

```
make mcprof.compile
```

The above command will generate the binary of the application. Next, to profile the application by *MCPProf*, run the following command:

```
make mcprof.execute
```

## 9 *MCPProf* Input Options

The complete list of the input options to *MCPProf* can be obtained by running the following command:

```
pin -t <Path to MCPProf dir>/obj-intel64/mcprof.so -h -- ls
```

Some of the important input options are detailed here.

**-RecordStack** [0/1, default 0] to tell *MCPProf* to include stack accesses or not.

**-TrackObjects** [0/1, default 0] to tell *MCPProf* if you want to track objects. If set 1, the calls to memory allocations functions (malloc, calloc, realloc, free, new, delete) will be instrumented to track the allocation and deallocation of objects at run-time in the application. The complete call-path to the allocation site, with source file-name and line-no information will be recorded as well and available in symbol.dat file.

**-Engine** [1/2/3, default 1] This selects the engine to be used in *MCPProf*. Based on the selected engine, the desired output is generated as below:

**Engine 1** provides the compute and memory intensive functions and objects (if TrackObjects is 1) in the application in the form of a flat profile and also generates an application callgraph. It also provides the data-communication information between functions. If TrackObjects is 1, the communication is also reported through objects in the application.

**Engine 2** provides information about a loop if it is independent. This feature is under development and not fully tested.

**Engine 3** reports the memory access information through functions and objects per call.

**-ShowUnknown** [0/1, default 0] Show Unknown function in the output graphs.

**-TrackStartStop** [0/1, default 0] Track start/stop markers in the code to start/stop profiling instead of starting from main().

**-TrackZones** [0/1, default 0] Track zone markers to profile per zone. See **markers.c** example available in tests directory. Note **markers.h** include. Also note how this test is compiled and executed in **makefile.rules**.

**-Threshold** [32-bit integer, default 0] to set the threshold value for the output generated by *MCPProf*. Communication below this threshold will not appear in the output graph.

## 10 *MCPProf* Generated Output

*MCPProf* generates various output files based on the selected engines. An important file which is generated independent of the selected engine is **symbols.dat**. This file contains the information about the function/object symbols tracked during the application execution. In case of objects, the allocation address, allocation size, call-path of the allocation is also reported.

## 10.1 Engine 1 Output

The output of Engine 1 are two text-files **execProfile.dat** in the current directory.

**execProfile.dat** file contains the information about the execution profile of the application. This profile lists the compute intensive functions in the application by reporting the percentage of dynamically executed instructions. An example output for the **vectOps** application is shown below:

%Exec.Instr.	Function Name
=====	=====
95	main
1	diffVecs
1	sumVecs
1	initVecs

**memProfile.dat** file contains the information about the memory accesses performed by functions/objects in the application. An example output for the **vectOps** application is shown below:

This table can be sorted by Total Accesses (-k2) by using bash command: tail -n +7 memProfile.dat   sort -k2 -gr						
Per Function Flat Memory Profile						
Function Name	=====	Accesses	=====	Allocation		
	Total	Reads	Writes	Path		
=====						
UnknownFtn	106270	79937	26333	NA:0		
main	78818	66750	12068	NA:0		
sumVecs	12000	8000	4000	tests/vectOps.c:64		
diffVecs	12000	8000	4000	tests/vectOps.c:65		
Per Object Flat Memory Profile						
Object Name	=====	Accesses	=====	====	====	====
	Total	Reads	Writes	Alloc %	Prod %	Path
=====						
Obj2	12000	8000	4000	100	100	tests/vectOps.c:50
Obj3	12000	8000	4000	100	100	tests/vectOps.c:53
Obj4	4004	4	4000	100	0.1	tests/vectOps.c:56
Obj5	4004	4	4000	100	0.1	tests/vectOps.c:59

The accesses reported against **UnknownFtn** are the accesses which cannot be associated to any function. For instance, the accesses before the **main** function is called.

In case of objects, Total, Read and Write memory accesses are reported in Column 2-4. Two other very useful efficiency factors are reported in column 5 and 6, under Efficiency. Column 5 (Alloc %) is the percentage of allocation efficiency. This indicated what percentage of allocated bytes are actually used for writing some data. Column 6 (Prod %) is the percentage of production efficiency. This is an indication of what percentage of data produced is actually consumed (read) later in the application before the object is freed. If production efficiency of an allocated object is low, then its not really utilized, which might mean that this data should not have been produced.

Last Column reports source-file and line-no of where the allocation has taken place in the application.

*MCPProf* generates call graph in simple text file, json format. The callgraph in json format can be visualized by the provided **callgraph2pdf.sh** script by the following command to generate the callgraph as shown in Figure 2.



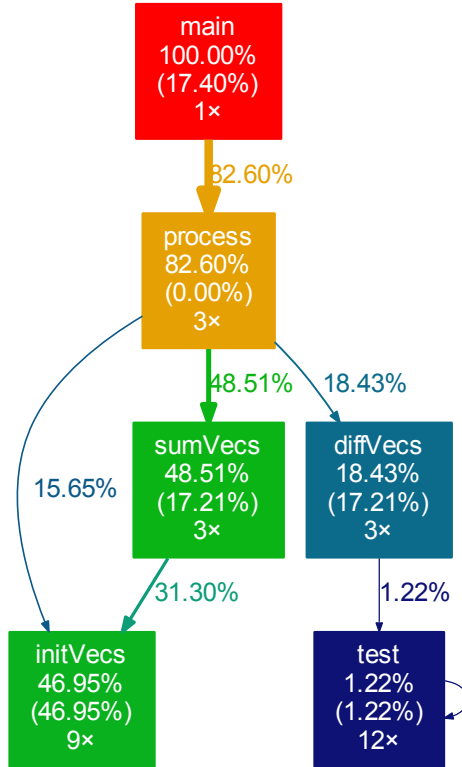


Figure 2: Call graph generated by *MCPProf*. The computational intensity is represented by the percentage of the dynamically executed instructions.

```
<path to mcprof dir>/scripts/callgraph2pdf.sh
```

*MCPProf* records data-communication among functions. This is reported as a data-communication in DOT format in the file **communication.dot** in the current directory. This file can be converted to pdf by the following command:

```
dot -Tpdf communication.dot -o communication.pdf
```

A script **dot2pdf.sh** is also provided in the **scripts** directory which can also be used to convert the graphs in *dot* format to *pdf* format. This script

also performs some extra tasks, for instance, remove unconnected nodes in the graph. This script can be used as:

```
<path to mcprof dir>/scripts/dot2pdf.sh
```

Figure 3 shows the data-communication graph generated by *MCPProf* with `TrackObjects` as 0.




Figure 3: Data-communication among functions in `vectOps` application as reported by *MCPProf*. The Grey ovals represent functions. The arcs represent the communication with the number on the arc representing the amount of data-communication in bytes. The number inside the ovals with % represent the percentage of the dynamically executed instructions. The second number is the total calls to this function.

Another important format of the reported information is the data-communication matrix. In this format, the communication is reported in the form of a matrix in a **matrix.dat** text file. This file can be converted to a graph by using `plotScript.sh` script available in **scripts** directory. This can be executed as:

```
<path to mcprof dir>/scripts/plotScript.sh
```

Figure 4 shows the resulting graph for `vectOps` application.

When `TrackObjects` is set to 1, then the object allocation (`malloc/new`) are also tracked and data-communication is reported through these objects. Figure 5 shows the same graph with `TrackObjects` as 1.

It is important to mention here that the names of the static objects are automatically detected by reading (ELF) header (will be available soon). However, the names of the dynamic objects are supplied by the user. *MCPProf* provides complete path to the allocation of an object in the source-code. So user has to manually look at the source-code to specify what should be the name of each object, in order to see the meaningful names of the objects. Hence, the user can insert these names directly in **communication.dot** file. A script will also be added soon to automate this process as much as possible.




Figure 4: Data-communication Matrix showing intensity of communication among functions in vectOps application as reported by *MCProf*.

## 10.2 Engine 2 Output

This engine simply provides the information if iterations of the selected loop are independent or not. In order to get an idea about how this can be achieved, reader is referred to see the details of the `tests/loopDepend.c` file and run the following to actually test it:

```
make loopDepend.test
```

## 10.3 Engine 3 Output

*MCProf* reports per call accesses in Engine 3 in `percallaccesses.dat` text file, as shown below:

```
Printing All Calls

Printing Calls to initVecs
Total Calls : 1
Call No : 0
Call Seq No : 1
Call Stack : UnknownFtn -> main -> initVecs
Writes to Object5 : 400
Writes to Object6 : 400

Printing Calls to sumVecs
Total Calls : 1
```




Figure 5: Data-communication among functions in vectOps application as reported by *MCPProf*. The tracked objects and communication through these objects is also shown. The Grey ovals represent functions. The white rectangles represent objects. The number inside the boxes represent the size of this object. The arcs represent the communication with the number on the arc representing the amount of data-communication in bytes. The number inside the ovals with % represent the percentage of the dynamically executed instructions. The second number is the total calls to this function.

```

Call No : 0
Call Seq No : 2
Call Stack : UnknownFtn -> main -> sumVecs
Reads from Object5 : 400
Reads from Object6 : 400
Writes to Object7 : 400

Printing Calls to diffVecs
Total Calls : 1
Call No : 0
Call Seq No : 3
Call Stack : UnknownFtn -> main -> diffVecs
Reads from UnknownFtn : 3115
Reads from Object5 : 400
Reads from Object6 : 400
Reads from Object7 : 4
Reads from Object8 : 4
Writes to UnknownFtn : 797
Writes to Object8 : 400

Printing Calls to main
Total Calls : 1
Call No : 0
Call Seq No : 0

```

```
Call Stack : UnknownFtn -> main
Reads from UnknownFtn : 69368
Writes to UnknownFtn : 3568
```

## 11 Frequently Encountered Problems

This section will cover some of the frequently encountered problems will setting-up/using *MCPProf*.

### 11.1 Pin: No such file or directory

This is because Pin needs 32 bit libraries (See [4]). On Ubuntu, this can be solved by:

```
sudo dpkg --add-architecture i386
sudo apt-get update
sudo apt-get install libc6:i386 libncurses5:i386 libstdc++6:i386
```

### 11.2 Pin Injection Mode Error

On some systems if Pin (parent) injection mode is not enabled by default then you see an error as shown below.

```
E:Attach to pid 13972 failed.
E: The Operating System configuration prevents Pin
E: from using the default (parent) injection mode.
E: To resolve, either execute the following (as root):
E: $ echo 0 > /proc/sys/kernel/yama/ptrace_scope
E: Or use the "-injection child" option.
E: For more information, regarding child injection,
E: see Injection section in the Pin User Manual.
```

Solution is also suggested in this message, which is to become root and enable this injection. This can be achieved by running the following two commands:

```
sudo -i
echo 0 > /proc/sys/kernel/yama/ptrace_scope
exit
```

## 12 Contact

In case you are interested in contributing to *MCPProf*, or you have suggestions for improvements, or you want to report a bug, contact:

- Imran Ashraf (I.Ashraf@TUDelft.nl )

## References

- [1] C.K. Luk and et al. Pin: Building Customized Program Analysis Tools with Dynamic Instrumentation. In *PLDI '05*, pages 190–200, New York, NY, USA, 2005. ACM.
- [2] Pin: A Dynamic Binary Instrumentation Tool. [www.pintool.org](http://www.pintool.org).
- [3] MCProf: Memory and Communication Profiler.
- [4] Libraries for Linux (Pin Requirement).